

Lab 3. Dynamic Memory Manager Module


System Programming

Juyoung Park
SNU TNET Lab.

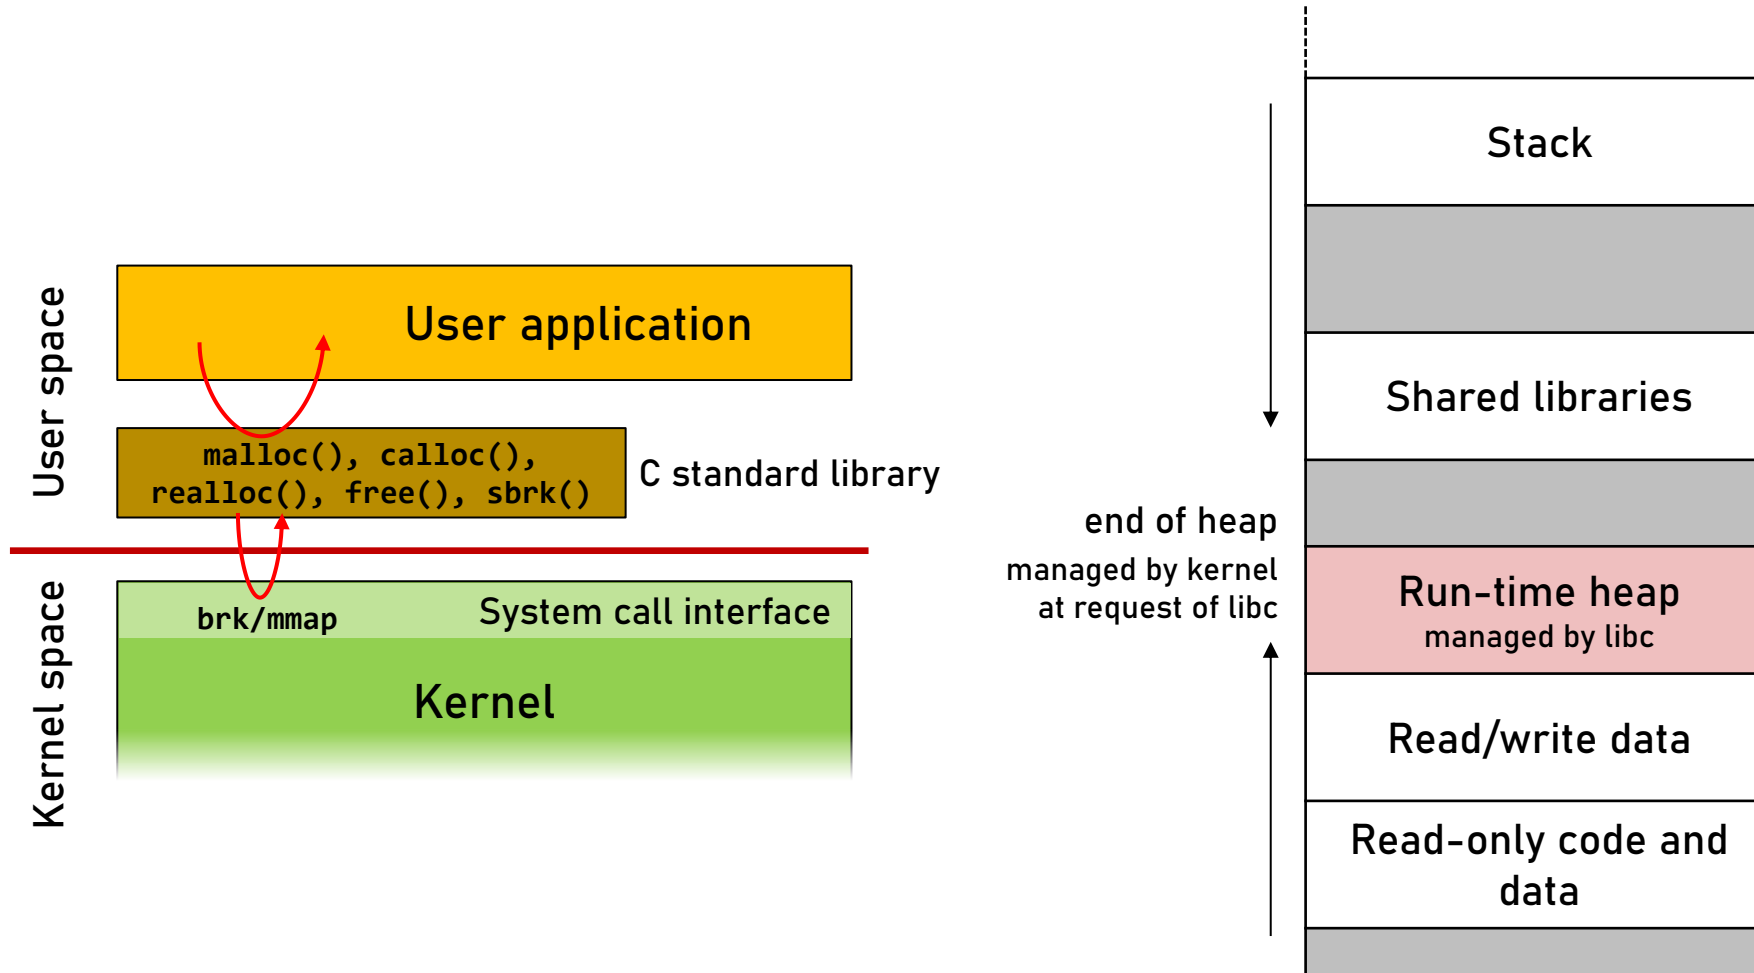
What Have You Learned

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
➡ free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```




What Have You Learned



What Have You Learned

- K&R Heap Manager



Dynamic Memory Manager Module

- Build a library that implements `malloc()` and `free()`
 - Without using GNU `malloc()`, `free()`, `calloc()`, or `realloc()`
- The code for your reference (`heapmgrgnu.c`, `heapmgrkr.c`) and baseline code (`heapmgrbase.c`) will be given
- Guidance on the assignment can also be found on `README.md`

Given Code - heapmgrgnu.c

- Implementation that simply calls the GNU malloc() and free()

```
/*-----*/
/* heapmgrgnu.c */
/* Author: Bob Dondero */
/* Using the GNU malloc() and free() */
/*-----*/

#include "heapmgr.h"
#include <stdlib.h>

/*-----*/

void *heapmgr_malloc(size_t ui_bytes)

/* Return a pointer to space for an object of size uiBytes. Return
   NULL if uiBytes is 0 or the request cannot be satisfied. The
   space is uninitialized. */

{
    return malloc(ui_bytes);
}

/*-----*/

void heapmgr_free(void *pv_bytes)

/* Deallocate the space pointed to by pvBytes. Do nothing if pvBytes
   is NULL. It is an unchecked runtime error for pvBytes to be a
   pointer to space that was not previously allocated by
   HeapMgr_malloc(). */

{
    free(pv_bytes);
}
```

Given Code - heapmgrkr.c

- Kernighan and Ritchie (K&R) implementation
 - With small modification for the sake of simplicity
 - `void *heapmgr_malloc(size_t nbytes)`
 - `void heapmgr_free(void *ap)`
 - `Header *morecore(unsigned int nu)`
- A circular, singly-linked list

Given Code - heapmgrbase.c

- Implements baseline code

- You can start the task with this code

```
void *heapmgr_malloc(size_t size)
```

```
void heapmgr_free(void *m)
```

```
int check_heap_validity(void)
```

- Validity check for entire data structures for chunks
- Other functions for implementation

- A non-circular, singly-linked list

Chunk Structure

```
/* chunkbase.h */
typedef struct Chunk *Chunk_T;

/* chunkbase.c */
struct Chunk {
    /* Pointer to the next chunk in the free chunk list */
    Chunk_T next;
    /* Capacity of a chunk (chunk units) */
    int units;
    /* CHUNK_FREE or CHUNK_IN_USE */
    int status;
};
```

Chunk and Block


- Chunk: base unit for allocate memory
- Block: set of contiguous chunks that store same data
- Free blocks are connected in a linked list
 - which is called a free list

Given Code - heapmgrbase.c

- `int check_heap_validity(void)`
 - Checks the validity of chunk data structures (chunk: a base unit for allocation)
 - Returns 1 on success or 0 (zero) on failure
- `assert(condition);`
 - If “condition” evaluates to false, the program will print an error message and terminate
- `heapmgrbase.c` calls `assert(check_heap_validity())`
 - At leading and trailing edges of `heapmgr_malloc()` and `heapmgr_free()`
 - Checks the integrity of the heap
 - If this `assert()` fails, it implies that something's wrong

Your to-do: Make free() faster


- Base requirement for assignment 3 (write code in heapmgr1.c)
- Implement a doubly-linked free list with the chunk data structure
 - Each chunk now contains a header and a footer (as described in lectures)
 - Chunk is a base unit (e.g., allocate memory in multiples of this unit)



- heapmgr1.c (with footer) should make free() faster
 - Without needing to find/maintain the “previous” node for inserting a freed memory into the free list (K&R)


Navigating Previous/Next Contiguous Block

- Start with the user's data portion of the block
- Go backwards to the head of the block
 - Easy, since you know the size of the header
- Go backwards to the footer of the previous block
 - Easy, since you know the size of the footer
- Go backwards to the header of the previous block
 - Easy, since you know the size from the footer



Navigating Previous/Next Free Block

- Start with the user's data portion of the block
- Go backwards to the head of the block
 - Easy, since you know the size of the header
- Go forwards to the footer of the block
 - Easy, since you know the block size from the header
- Go backwards to the previous free block
 - Easy, since you have the previous free pointer




Strategy for Writing heapmgr1.c

- You can start with heapmgrbase.c and enhance it
 - Or you can ignore heapmgrbase.c and do your own way
- You can implement either circular or non-circular list

Extra-credit: Make malloc() faster

- heapmgr1.c shows poor worst-case behavior, so try to enhance it
 - Finding the free block will traverse the list in heapmgr1.c
- Use multiple doubly-linked lists, alias bins (as described in lectures)



Strategy for Binning

- How to set the range of sizes covered by bin?
 - Fixed: 1-10, 11-20, 21-30, ...
 - Exponential: 1-2, 3-4, 5-8, 9-16, ...
- How to handle the newly allocated large memory?
 - Split into small chunks in advance
 - Allocate in large bin and wait for split in malloc()
- How do you coalesce when free() from binning?
 - Check after coalesce and move to proper bin
 - Move to final bin first and coalesce

chunk.c and chunk.h

- heapmgrbase.c uses chunkbase.c and chunkbase.h to use 'struct Chunk' and functions
- If you use heapmgrbase.c as skeleton code, you can use chunkbase.c and chunkbase.h for chunk.c and chunk.h
- If you do not use heapmgrbase.c, you have two choices:
 - 1. Define useful structures and functions in chunk.c and chunk.h and use them
 - 2. Just leave chunk.c and chunk.h as is and not use them

Useful Functions in chunkbase.c & chunkbase.h

- `chunk_get_status(Chunk_T c)`
 - Returns a chunk's status
- `chunk_set_status(Chunk_T c, int status)`
 - Set the status of the chunk
- `chunk_get_units(Chunk_T c)`
 - Returns the size of a chunk
- `chunk_set_units(Chunk_T c, int units)`
 - Sets the current size in 'units' of 'c'

Useful Functions in chunkbase.c & chunkbase.h

- `chunk_get_next_free_chunk(Chunk_T c)`
 - Returns the next free chunk in free chunk list
- `chunk_set_next_free_chunk(Chunk_T c, Chunk_T next)`
 - Sets the next free chunk of 'c' to 'next'
- `chunk_get_next_adjacent(Chunk_T c, void *start, void *end)`
 - Returns the next adjacent chunk to 'c' in memory space
- `chunk_is_valid(Chunk_T c, void *start, void *end);`
 - Checks the validity of a chunk

Memory Utilization

- Implement the strategies for good memory utilization
 - All techniques learned in class
 - Check blocks in free list before allocate new memory
 - Divide the free block if free block is bigger then requested
 - Check lower/upper neighbor and coalesce
- If you ignore memory utilization, you won't get points no matter how fast your implementation is

How to Test Your Code

- To test your heapmgr implementations:
 - `$ gcc800 -std=gnu99 testheapmgr.c heapmgr1.c chunk.c -o testheapmgr1`
 - `$ gcc800 -std=gnu99 testheapmgr.c heapmgr2.c chunk.c -o testheapmgr2`
- To collect timing statistics:
 - `$ gcc800 -O3 -D NDEBUG -std=gnu99 testheapmgr.c heapmgrgnu.c -o testheapmgrgnu`
 - `$ gcc800 -O3 -D NDEBUG -std=gnu99 testheapmgr.c heapmgrkr.c -o testheapmgrkr`
 - `$ gcc800 -O3 -D NDEBUG -std=gnu99 testheapmgr.c heapmgrbase.c chunkbase.c -o testheapmgrbase`
 - `$ gcc800 -O3 -D NDEBUG -std=gnu99 testheapmgr.c heapmgr1.c chunk.c -o testheapmgr1`
 - `$ gcc800 -O3 -D NDEBUG -std=gnu99 testheapmgr.c heapmgr2.c chunk.c -o testheapmgr2`
- Don't forget **-std=gnu99** ; otherwise you'll get error while compiling

How to Test Your Code

- You can also use Makefile to build executable files

make +	commands to be executed
test1	gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o test/testheapmgr1
test2	gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr2.c src/chunk.c -o test/testheapmgr2
testall	gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o test/testheapmgr1 gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr2.c src/chunk.c -o test/testheapmgr2
timegnu	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrgnu.c -o test/testheapmgrgnu
timekr	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrkr.c -o test/testheapmgrkr
timebase	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrbase.c reference/chunkbase.c -o test/testheapmgrbase
time1	gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o test/testheapmgr1
time2	gcc800 -std=gnu99 test/testheapmgr.c src/heapmgr2.c src/chunk.c -o test/testheapmgr2
time1all	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrgnu.c -o test/testheapmgrgnu gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrkr.c -o test/testheapmgrkr gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrbase.c reference/chunkbase.c -o test/testheapmgrbase gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o test/testheapmgr1
time2all	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrgnu.c -o test/testheapmgrgnu gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrkr.c -o test/testheapmgrkr gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrbase.c reference/chunkbase.c -o test/testheapmgrbase gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o test/testheapmgr1 gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c src/heapmgr2.c src/chunk.c -o test/testheapmgr2
all (same as time2all)	gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrgnu.c -o testheapmgrgnu gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrkr.c -o testheapmgrkr gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c reference/heapmgrbase.c reference/chunkbase.c -o test/testheapmgrbase gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c src/heapmgr1.c src/chunk.c -o testheapmgr1 gcc800 -O3 -D NDEBUG -std=gnu99 test/testheapmgr.c src/heapmgr2.c src/chunk.c -o test/testheapmgr2
clean	rm -f test/testheapmgrgnu test/testheapmgrkr test/testheapmgrbase test/testheapmgr1 test/testheapmgr2

How to Test Your Code

- If you want to use Makefile, move the codes you implemented (chunk.c, chunk.h, heapmgr1.c, heapmgr2.c) in src folder
- You don't need to move the existing codes
- Executable files will be created in the test folder

```
assignment3
|- reference
|  |- chunkbase.c / chunkbase.h
|  |- heapmgr.h
|  |- heapmgrgnu.c / heapmgrkr.c
|  `-- heapmgrbase.c
|- src
|  |- chunk.c / chunk.h
|  `-- heapmgr1.c / heapmgr2.c
|- test
|  |- heapmgr.h
|  `-- testheapmgr.c
`--Makefile
```


How to Test Your Code

- Bash shell scripts (testheap1 & testheap2) will be provided
- testheap1 runs testheapmgr.c to test four cases (heapmgrgnu.c, heapmgrkr.c, heapmgrbase.c, heapmgr1.c) and reports timing and memory usage statistics

```
#!/bin/bash

#####
# testheap1 tests four HeapMgr implementations.
# Executable files named testheapmgrgnu, testheapmgrkr, testheapmgrbase,
# and testheapmgr1 must exist before executing this script.
# To execute the script, simply type testheap1.
#####

echo "      Executable      Test   Count   Size   Time      Mem"
./testheapimp ./testheapmgrgnu
./testheapimp ./testheapmgrkr
./testheapimp ./testheapmgrbase
./testheapimp ./testheapmgr1
# ./testheapimp ./testheapmgr2
```

How to Test Your Code

- testheap2 runs testheapmgr.c to test five cases (heapmgrgnu.c, heapmgrkr.c, heapmgrbase.c, heapmgr1.c, and heapmgr2.c) and reports timing and memory usage statistics

```
#!/bin/bash

#####
# testheap2 tests five HeapMgr implementations.
# Executable files named testheapmgrgnu, testheapmgrkr, testheapmgrbase,
# testheapmgr1 and testheapmgr2 must exist before executing this script.
# To execute the script, simply type testhea2p.
#####

echo "      Executable      Test  Count  Size  Time      Mem"
./testheapimp ./testheapmgrgnu
./testheapimp ./testheapmgrkr
./testheapimp ./testheapmgrbase
./testheapimp ./testheapmgr1
./testheapimp ./testheapmgr2
```

How to Test Your Code

Argument	Test Performed
LIFO_fixed	LIFO with fixed size chunks
FIFO_fixed	FIFO with fixed size chunks
LIFO_random	LIFO with random size chunks
FIFO_random	FIFO with random size chunks
random_fixed	Random order with fixed size chunks
random_random	Random order with random size chunks
worst	Worst case order for a heap manager implemented using a single linked list

How to Test Your Code

- `$./testheap1 (or ./testheap2)`
 - Provides timing and memory usage statistics for all codes
- `$./testheapmgr1 LIFO_fixed 100 1000`
 - Perform a LIFO_fixed test with testheapmgr1
 - Run heapmgr_malloc() and heapmgr_free() 100 times
 - The (maximum) size of each memory chunk is 1000 bytes
- testheap1 and testheap2 is in test folder

How to Test Your Code

- Set the product of the **number of calls** and **size in bytes** to less than 10^9
 - \$./testheapmgr1 LIFO_fixed **100000 100000 (X)**
- In all tests evaluating the implementation on the Bacchus machine, (number of cells) x (size in bytes) $\leq 10^9$ is guaranteed
- Only tests for final check should be performed on the Bacchus machine
 - Otherwise, test in your local machine

Content of readme file

- Your name and student ID
- Result of `./testheap1` or `./testheap2` (paste the output of the `testheap1` or `testheap2` script)
- (Optionally) An indication of how much time you spent doing the assignment
- (Optionally) Your assessment of the assignment
- (Optionally) Any information that will help us to grade your work in the most favorable light

How to Submit?

- Make a directory
 - `$ mkdir 202400000_assign3`
- Move your code and readme file there
 - `$ mv heapmgr1.c (heapmgr2.c) chunk.c chunk.h readme 202400000_assign3`
- Make a gzipped tar file for submission
 - `$ tar zcf 202400000_assign3.tar.gz 202400000_assign3`
- You **must** submit chunk.c and chunk.h even if you did not use them in your implementation
 - In this case, submit chunk.c and chunk.h given in the assignment without modification

How to Submit?

Structure of directory:

YourID_assign3 (don't use dash)

|-heapmgr1.c

|-heapmgr2.c (optional)

|-chunk.c

|-chunk.h

`-readme (don't use extension such as .txt, .md, ...)

Example:

202400000_assign3

|-heapmgr1.c

|-chunk.c

|-chunk.h

`-readme

- Please set **files and directory's names** to match the examples above
 - Don't use any extension for readme file
 - Don't use dash for submit file
- Structure files and directories as shown above, then proceed with **compression**
- **Deadline: ~11.1(Fri) 21:00**
 - **0 points if deadline is missed**

Grading - heapmgr1.c

- Submit format (12 / 100)
 - readme, files with proper names
- Evaluation from the user' viewpoint: function correctness (78 / 100)
 - heapmgr_malloc(), heapmgr_free() are well-designed (include validity check)
 - time consumption: faster than heapmgrkr.c, heapmgrbase.c (except worst)
- Evaluation from the programmer's viewpoint (10 / 100)
 - Clarity (names, comments, line lengths, indentation, etc.)
 - Parameter validation using assert()

Grading - heapmgr2.c (Extra Credit)

- Evaluation from the user' viewpoint: function correctness (+30%)
 - heapmgr_malloc(), heapmgr_free() are well-designed (using bins) (include validity check)
 - time consumption: faster than heapmgrkr.c, heapmgrbase.c and heapmgr1.c
- Extra credit
 - Up to 30% of the scores you earn for implementing heapmgr1.c
- You will not get the 30% extra credit just by submitting heapmgr2
 - If you do not get a perfect score on heapmgr2, extra credit you receive will be reduced