

Lab 4. A Unix Shell

SNU System Programming Assignment

Jongki Park
SNU TNET Lab.

What You Should Do

1. Implement basic shell functionality

- Utilize parsed command line
- External command execution

2. Support for redirection

- Standard input redirection
- Standard output redirection

3. Support for pipe

- Single pipe support
- Multiple pipe support

Extra Credit for Background Process Support

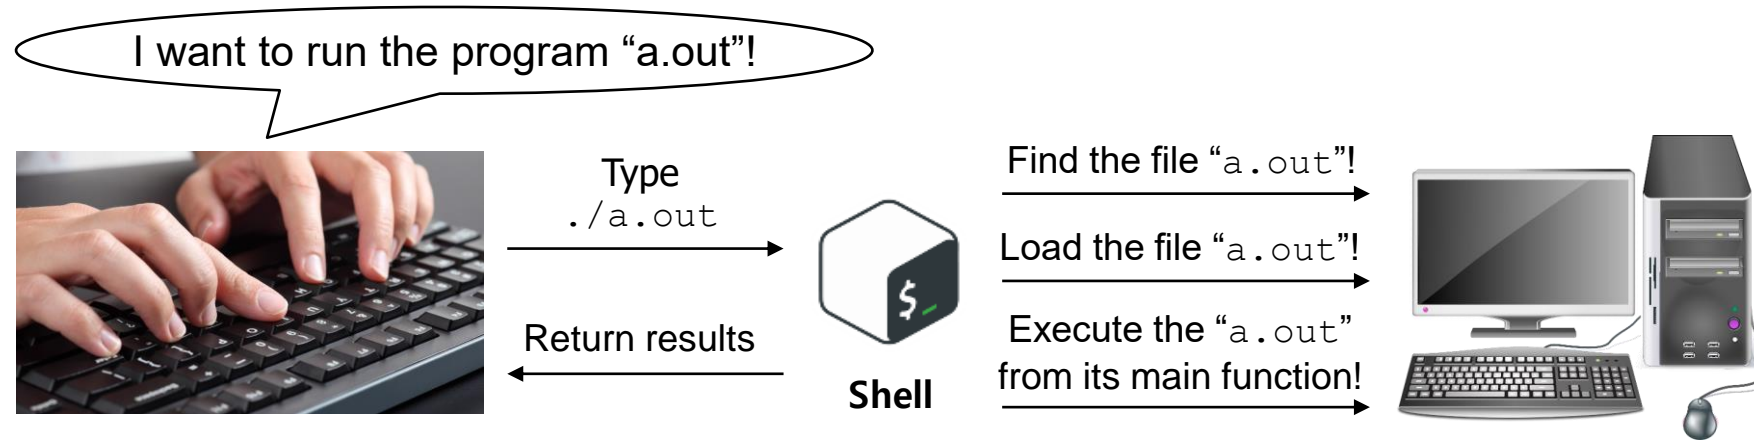
4. Support for background process

- Running a process in the background mode
- Add job(s) to a background job list
- Handle the process when the background process terminates

Basic Shell Functions

What Is a Shell?

- Shell: a program that executes commands typed by a user



A shell =

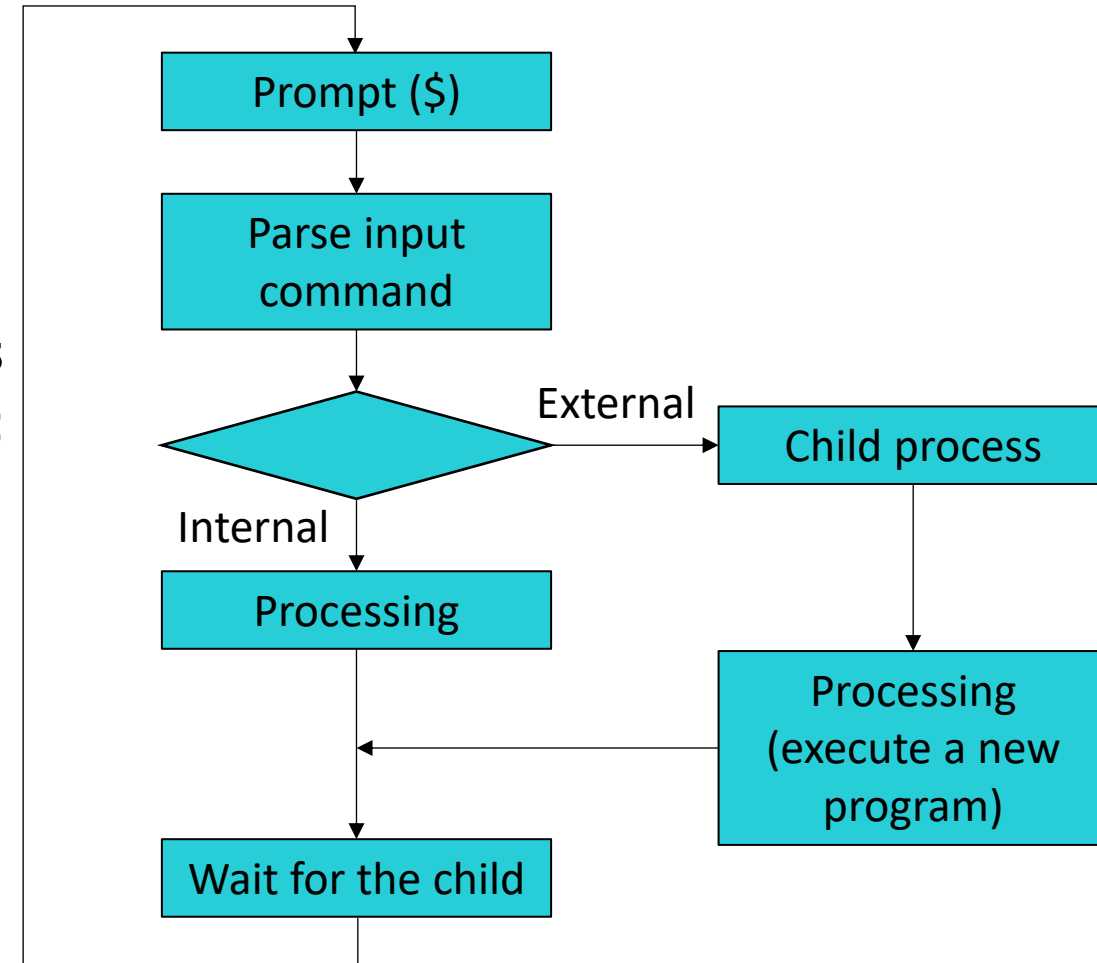
```
while(1) {
```

- Receives a input command string from the user
- Interprets the command
- Executes the command requested by user

```
}
```

Basic Shell Functionality

- Command interpreter
 - Execute commands requested by the user
- Basic logic
 - Display the prompt & parse input commands
 - External commands: fork and exec at child process
 - Internal commands: process without fork and exec
 - Parent process (Shell) waits for the child to end



System Calls for Basic Shell Functions

- `fork()` creates a child process that duplicates the parent
 - Manipulating fds after `fork()` does not affect the fds of the other process
- `execvp()` finds, loads, and executes an external command in current process
 - `Execvp()` automatically searches directories in `$PATH` (env. var.) for the executable binary
 - The first parameter is the binary name without any path (`"ls"` not `"/usr/bin/ls"`)
- `waitpid()` is called by the parent process to reap a terminated child process
 - Parent blocks until a state of a child is changed to `EXITED`
 - `SIGCHLD` signal is sent to the parent when its child terminates
 - `SIGCHLD` signal handler of parent calls `waitpid()` to reap the terminated child

A Minimal Shell for Executing a Single Command

- Simple shell code
- Notice
 - Return value/error checking is ***mandatory*** when doing the assignment to ensure proper error handling
 - The code here is simplified just for easy understanding

```
fgets(command, sizeof(command), stdin);  
command[strlen(command)-1] = 0;
```

① **Get the command**

```
token = strtok(command, " ");  
if (token == NULL) {exit(-1);}
```

② **Parse(Tokenize) command**

```
arguments[0] = token;  
for (i = 1; i < 10; i++) {  
    token = strtok(NULL, " ");  
    if (token == NULL)  
        break;  
    arguments[i] = token;  
}  
arguments[i] = NULL;
```

③ **Compose argument array**

```
pid = fork();
```

④ **Fork child process**

```
if (pid == 0){  
    execvp(arguments[0], arguments);  
}
```

⑤ **Child process:**
changes its execution
image to arguments[0]

```
/* Parent process */  
pid = wait(NULL);
```

⑥ **Parent process :**
waits for exit of the child

Your T0-D0s for snush: Basic Shell Functions

Functions for Basic Shell Functions

- No need to implement handling internal (built-in) commands
- Two built-in commands are provided
 - `cd` command
 - Moves to the specified directory
 - If no argument is provided, use the env. variable (`HOME`) to navigate to the home directory
 - `exit` command to exit the shell `snush`
- You need to implement `fork_exec()`
- Important system calls to utilize
 - `fork()`, `execvp()`, `waitpid()`
- Important `snush` functions to utilize()
 - `build_command()`: a wrapper for `build_command_partial()`
 - Uses a full range of token indices
 - Do not modify `build_command()` & `build_command_partial()`

Implementing Basic Shell Functions

- Create a child process
- Convert the tokenized input into an argument array for an external command
- Do appropriate actions for supporting I/O redirection or a pipe
- Execute the command in the child process
- Make sure the parent process waits properly
- The return value of `fork_exec()` should be the child's `PID`

Support for Redirection


System Calls for Redirection

- Unix allows programmatic way to redirect `stdin`, `stdout` or `stderr`
- How?
 - Use `open()` or `create()` to open a file for reading or writing, creating the file
 - Use `close()` to close unused file descriptors
 - Use `dup2()` to duplicate the file descriptor to any target descriptor, including `stdin`, `stdout` or `stderr`
- For this assignment, support only `stdin/stdout` redirection (no `stderr`)

Simple Redirection Trace (1)

- Simple code for implementing “someprogram > somefile”
- A new child has the identical file descriptor table as parent's
 - The first three file descriptors of a child point to the same terminal as parent's

```
pid = fork();  
if (pid == 0){  
    /* Child process */  
    fd = creat("somefile", 0640);  
    dup2(fd, 1);  
    close(fd);  
    execvp("someprogram", program_args);  
    fprintf(stderr, "exec failed\n");  
    exit(EXIT_FAILURE);  
}  
  
/* Parent process */  
pid = wait(NULL);
```




Simple Redirection Trace (2)

- Child process creates `somefile`

```
pid = fork();

if (pid == 0){
    /* Child process */
    fd = creat("somefile", 0640);
    dup2(fd, 1);
    close(fd);
    execvp("someprogram", program_args);
    fprintf(stderr, "exec failed\n");
    exit(EXIT_FAILURE);
}

/* Parent process */
pid = wait(NULL);
```




Simple Redirection Trace (3)

- `fd` and file descriptor number 1 can be used interchangeably after the successful call of `dup2()`

```
pid = fork();

if (pid == 0){
    /* Child process */
    fd = creat("somefile", 0640);
    dup2(fd, 1);
    close(fd);
    execvp("someprogram", program_args);
    fprintf(stderr, "exec failed\n");
    exit(EXIT_FAILURE);
}

/* Parent process */
pid = wait(NULL);
```




Simple Redirection Trace (4)

- Closes `fd` (index 3 at file descriptor table)

```
pid = fork();

if (pid == 0){
    /* Child process */
    fd = creat("somefile", 0640);
    dup2(fd, 1);
    close(fd);
    execvp("someprogram", program_args);
    fprintf(stderr, "exec failed\n");
    exit(EXIT_FAILURE);
}

/* Parent process */
pid = wait(NULL);
```




Simple Redirection Trace (5)

- someprogram **executes with** stdout **redirected to** somefile

```
pid = fork();

if (pid == 0){
    /* Child process */
    fd = creat("somefile", 0640);
    dup2(fd, 1);
    close(fd);
    execvp("someprogram", program_args);
    fprintf(stderr, "exec failed\n");
    exit(EXIT_FAILURE);
}

/* Parent process */
pid = wait(NULL);
```



Simple Redirection Trace (6)

- `someprogram` exits
- Parent process returns from `wait()` and proceeds


```
pid = fork();

if (pid == 0){
    /* Child process */
    fd = creat("somefile", 0640);
    dup2(fd, 1);
    close(fd);
    execvp("someprogram", "program_args");
    fprintf(stderr, "exec failed\n");
    exit(EXIT_FAILURE);
}

/* Parent process */
pid = wait(NULL);
```

Parent process

File
descriptor
table



Your T0-D0s for snush: Redirection

Functions for Redirection

- You need to implement `redout_handler()`
- Sample implementation of `redin_handler()` is provided in the skeleton code
- Important system calls to utilize
 - `open()`
 - `close()`
 - `dup2()`

Implementing Redirection in snush

- Redirect its input and output to a specified file before the command is executed
- `stdin` redirection: opens a file (w/ read-only perm.), replaces `stdin` with the fd
- `stdout` redirection: replaces `stdout` with the redirected fd
 - Set appropriate permissions for the output file
- No need to support `stderr` redirection

Support for Pipe

System calls for Pipe

- A pipe allows two processes on the same machine to exchange data
 - `stdout` of the previous program flows into `stdin` of the next program
- How?
 - Use `pipe()` to create a pair of fds (2 fds) that are connected to each other
 - `pipe()` creates a unidirectional communication line (one for read-only, the other for write-only)
 - Use `fork()` to create a child process that shares the two fds with the parent
 - Use `dup2()` to duplicate any fd, including `stdin`, `stdout` or `stderr`
- Pipe descriptors are file descriptors in UNIX


Example Use of Pipes

- The file "student_ids.txt" contains many student entries
- Find all students with the last name "Park"
 - `$ grep Park student_ids.txt`
- List the results in order of student id
 - `$ grep Park student_ids.txt | sort -n`

```
jongki@sp04:~$ cat student_ids.txt
202400001 Junghan Yoon
202400003 Seongjong Bae
202400009 Jongki Park
202400005 Juyoung Park
jongki@sp04:~$ grep Park student_ids.txt
202400009 Jongki Park
202400005 Juyoung Park
jongki@sp04:~$ grep Park student_ids.txt | sort -n
202400005 Juyoung Park
202400009 Jongki Park
jongki@sp04:~$
```

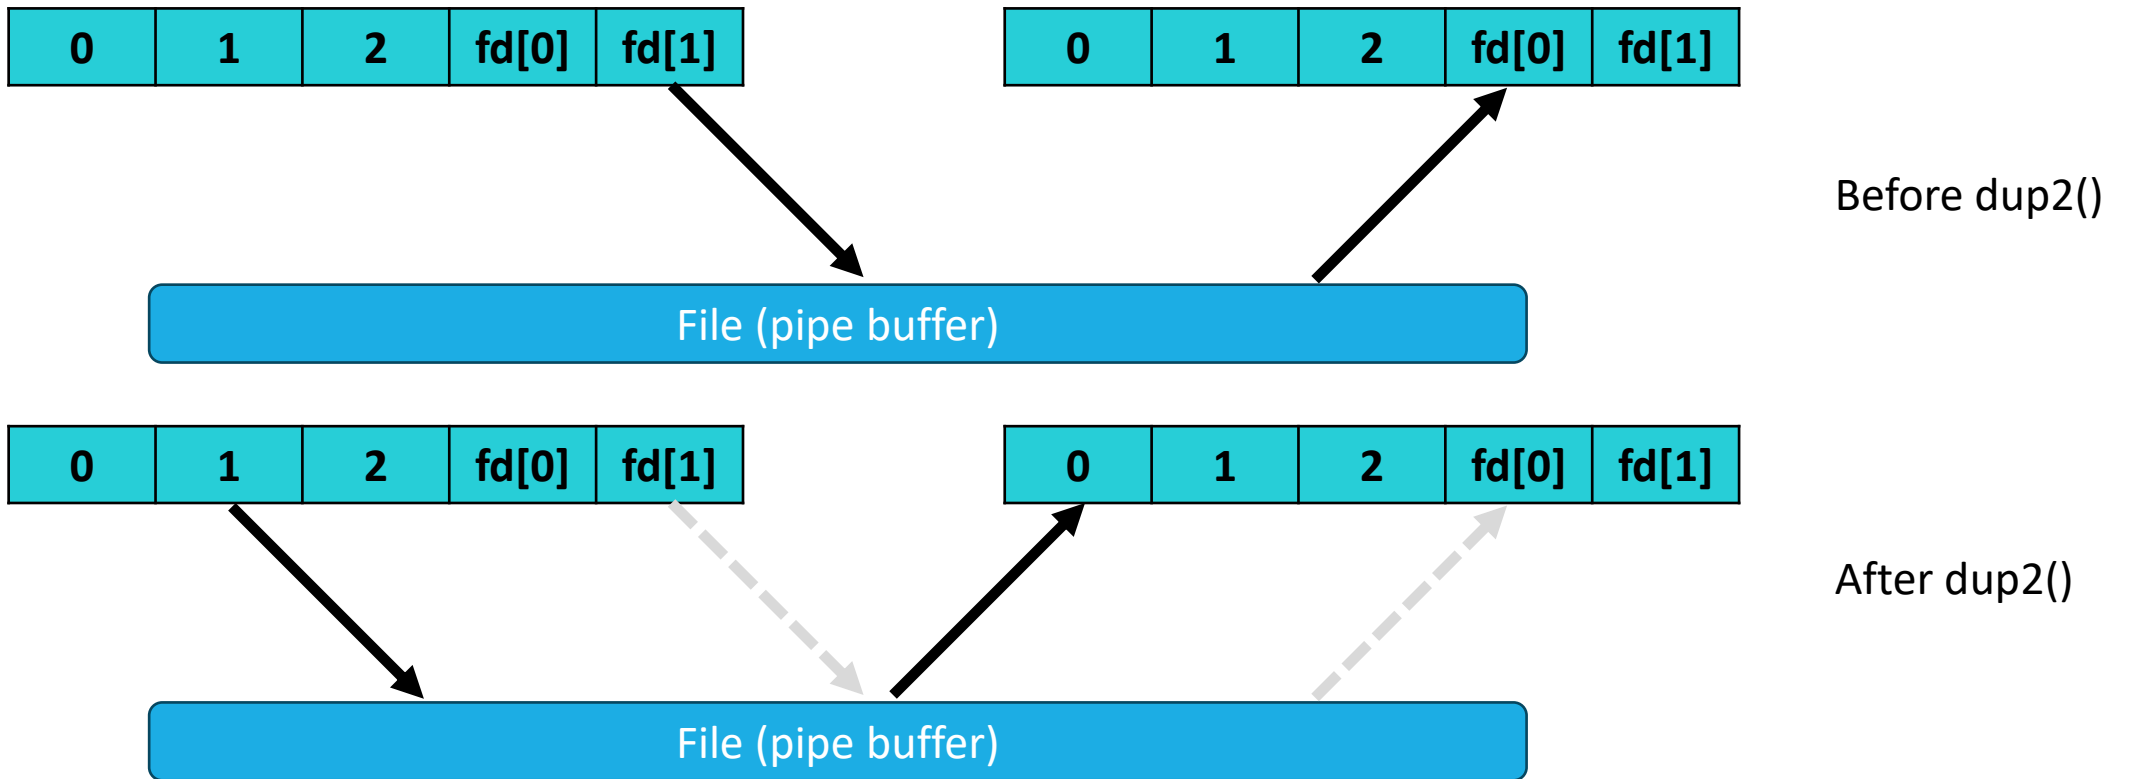
Unidirectional Communication

- `pipe()` creates a communication line (`fd[0]` for read-only, `fd[1]` for write-only)
- `pipe()`, `fork()`
 - Executing `fork()` after `pipe()` inherits the same file descriptors pointing to the same file



Duplicate to Other File Descriptors

- `dup2 ()`
 - Use `dup2 ()` to duplicate `stdin`, `stdout`, or other file descriptors to the file descriptors created by `pipe ()`
 - The output of the process can be passed to the input of the next process



Simple Pipe example

- Example code
 - Inter Process Communication between the child and the parent


```
int pid, p[2];
if (pipe(p) == -1)
    exit(1);

pid = fork();
if (pid == 0) {
    close(p[1]);
    dup2(p[0], 0);
    close(p[0]);
    ... read from stdin ...
}
else {
    close(p[0]);
    dup2(p[1], 1);
    close(p[1]);
    ... write to stdout ...
    wait(&status);
}
```

① Create pipe read: fd[0],
write: fd[1]

② Child makes stdin (0) the
read side of the pipe

③ Parent makes stdout (1) the
write side of the pipe



Your T0-D0 for snush: Pipe

Functions for Pipe

- You need to implement `iter_pipe_fork_exec()`
- Important system calls to utilize
 - `fork()`
 - `dup2()`
 - `close()`
 - `execvp()`
 - `waitpid()`
- Important `snush` functions to utilize (Do not modify the functions listed below)
 - `dynarray_get()`: get tokens of a specific index, where index starts at 0
 - `dynarray_get_length()`: get total number of tokens
 - `build_command_partial()`: construct the command array passed as argument to `execvp()`

Implementing Pipe in snush

- Decompose the token array based on the pipe symbol (`|`)
 - Note that the pipe symbol cannot be preceded by an output (`>`) or input (`<`) redirection operator
- Set up pipes
- Create child processes and run each command on each child
 - Create a child process for each command in the pipe
 - Convert the tokenized input to an argument array for external command execution
 - Execute a specified command at the child process
- The output of the previous process becomes the input of the next process
- Make sure the parent process waits properly
- `iter_pipe_fork_exec()` should return the `PID` of the first child

Support for Background Execution

Example for Background Process Execution

- Run tasks in the background and work on other commands simultaneously
- Useful for long-running processes (e.g., file downloads, backups, or large data processing)
- Child sends `SIGCHLD` signal to parent process when the child process exits

Simple Background Execution

- If parent process wants to wait for the child to finish, call `waitpid()`
- Otherwise, do not wait for the child to end its execution

```
pid = fork();

if (pid < 0) {
    perror("fork failed");
    exit(EXIT_FAILURE);
} else if (pid == 0) {
    execl("/bin/sleep", "sleep", "20", NULL);
    // execlp("sleep", "sleep", "20", NULL);

    perror("execlp failed");
    exit(EXIT_FAILURE);
} else {
    if (!background) {
        waitpid(pid, &cstatus, 0);
    } else {
        printf("Background: child process (PID: %d)\n", pid);
    }
    exit(EXIT_SUCCESS);
}
```

- ① Create a child process using `fork()`
- ② If it is a child process, change its execution image to another
- ③ If it is a parent process, decide whether to wait or not
- ④ If it is needs a background execution, a parent process do not wait for the child to exit

Your T0-D0 for snush: Background Execution

Functions for Background Execution

- You need to implement `sigzombie_handler()`
- Important system calls to utilize
 - `waitpid()`: set the third argument of `waitpid()`, the options field, to `WNOHANG` in `sigzombie_handler()`
- Remove the process ID from the list of background jobs
 - `bg_array[]`: global variable for the list of background jobs

Implementing Background Execution in snush

- Background processes are added to the background process array by the parent process
- The parent process does not wait for these background processes to complete
- Child sends `SIGCHLD` signal to parent process when the child process exits
- Signal handler function for `SIGCHLD` receives and removes each background process from the array upon their exit

Final Guidelines for snush

Input Command Line Usage Rules

- It is allowed to use symbols and commands without spaces in between
- The background execution symbol `&` cannot appear in the middle of the command line
- A redirection after a pipe is not valid
- There must be a target after a redirection symbol
- Nested redirections are valid only when input redirection `<` comes first, followed by output redirection `>`
- A pipe symbol cannot precede an input redirection symbol

Notice

- Do not change the name and the prototype of the skeleton code
- What to submit
 - Directory name should be `202400000_assign4`
 - Place the entire source codes in a single directory `202400000_assign4`
 - Make a gzipped tar file for submission

Deadline

- **Deadline: ~ 2024. 11. 28 21:00**
 - 0 Points if deadline is missed
 - 0 Points for copying
- **Contact**
 - Lab 4 TA e-mail: jkipark@snu.ac.kr
 - TA mailing list: snu-sysp@googlegroups.com
- Directory name should be `202400000_assign4`
- Place the entire source codes in a single directory `202400000_assign4`
- Make a gzipped tar file for submission

Test cases for self-checking

- All test cases for scoring will be provided, but scores may vary for each test case
- We provide test cases in the `self_check` directory:
 - Redirection input test
 - Redirection output test
 - Finding execution file test
 - Single pipe test
 - Multiple pipe test
 - Multiple pipe and output redirection test
 - Slow pipe test
 - Multiple slow pipe and output redirection test
 - Interrupt on single process test
 - Interrupt on multiple process test
 - Background execution test (Extra credit)

Appendix

snush Function for Parsing Input Command Line

- Take advantage of the command line parser that is already implemented
 - `lex_line()`: gets a command from stdin, parses, and store the array of token pointers
 - `syntax_check()`: checks if generated tokens are in an executable token format
 - Do not modify `lex_line()` and `syntax_check()`
- Do not modify these functions

The Result of Parsing in snush

- Set env. variable (DEBUG) before running your shell to see the tokenized output
 - `$ export DEBUG=1`
 - `$ unset DEBUG` (to turn it off)
- Example of how the `lex_line()` function works
 - `$ ls | sort > output &`
- Syntax check
 - Function `syntax_check()`
 - Performs an analysis to see if the tokenization with the `lex_line()` is in the executable form

```
% ls | sort >output &  
[0] TOKEN_WORD("ls")  
[1] TOKEN_PIPE(|)  
[2] TOKEN_WORD("sort")  
[3] TOKEN_REDIRECTION_OUT(>)  
[4] TOKEN_WORD("output")  
[5] TOKEN_BACKGROUND(&)
```

Token
Type : TOKEN_WORD
Value : ls

Token
Type : TOKEN_PIPE
Value : \|0

Token
Type : TOKEN_WORD
Value : sort

Token
Type : TOKEN_REDIRECTION_OUT
Value : \|0

Token
Type : TOKEN_WORD
Value : output

Token
Type : TOKEN_BG
Value : \|0

snush Functions for Checking the Presence of Pipe or Background

- Pipe / background command check
 - `count_pipe()`: counts the number of pipes in the command
 - `check_bg()`: checks if background process control operator `'&'` exists
- Do not modify these functions

END