

Lab 5. Simple KVS

SNU System Programming Assignment

Junghan Yoon
SNU TNET Lab.

What You Should Do

1. Implement basic server/client

- Use socket APIs
- Use multiple threads
- Use skvs helper function

2. Implement a global hash table and rwlock

- Multiple threads may access at the same time
- Global hash table uses rwlock library
- rwlock should support multiple concurrent readers when there is no writer

▪ Reference for socket programming in C:

- <https://beej.us/guide/bgnet/>

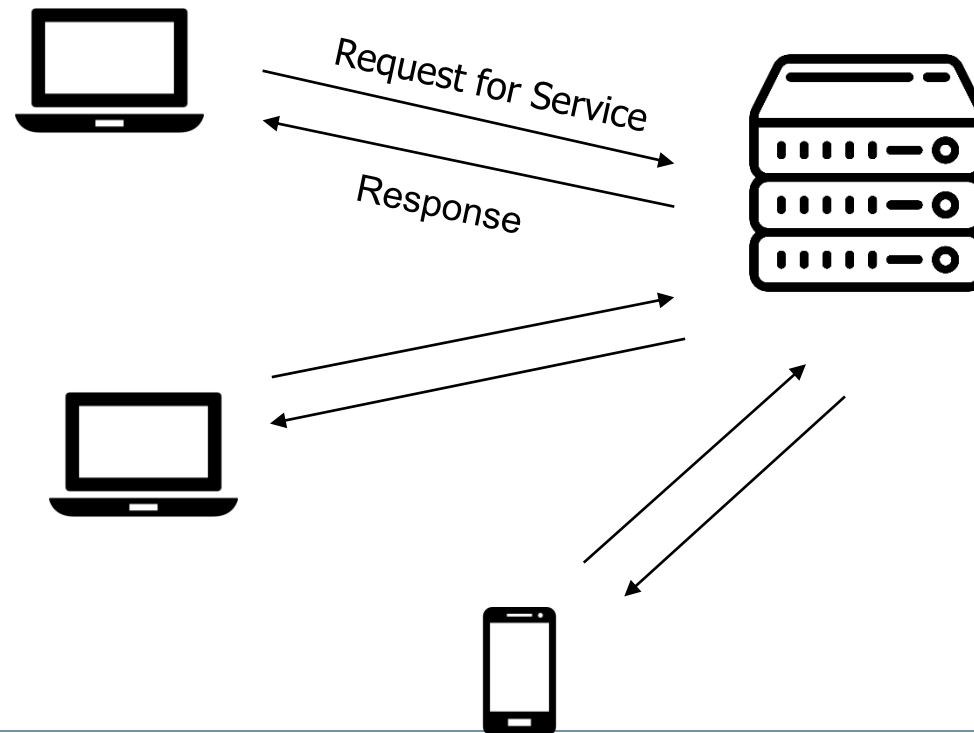
▪ Reference for rwlock:

- https://docs.oracle.com/cd/E37838_01/html/E61057/sync-124.html
- <https://www.ibm.com/docs/en/aix/7.3?topic=p-pthread-rwlock-rdlock-pthread-rwlock-tryrdlock-subroutines>

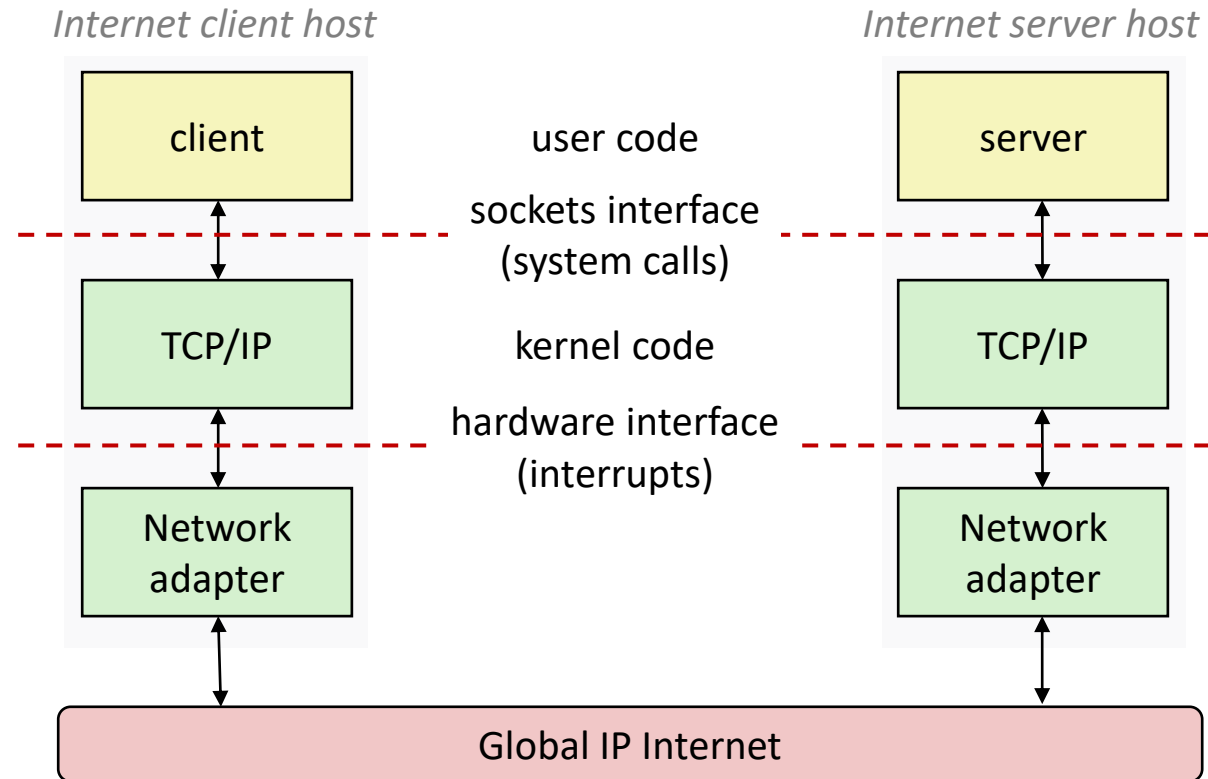
Background 1: Basic Server/Client

What is the Client-Server Model?

- **Definition:** A distributed computing architecture where a **client** requests services, and a **server** provides them.
- **Structure:** Typically involves multiple clients connecting to a centralized server.



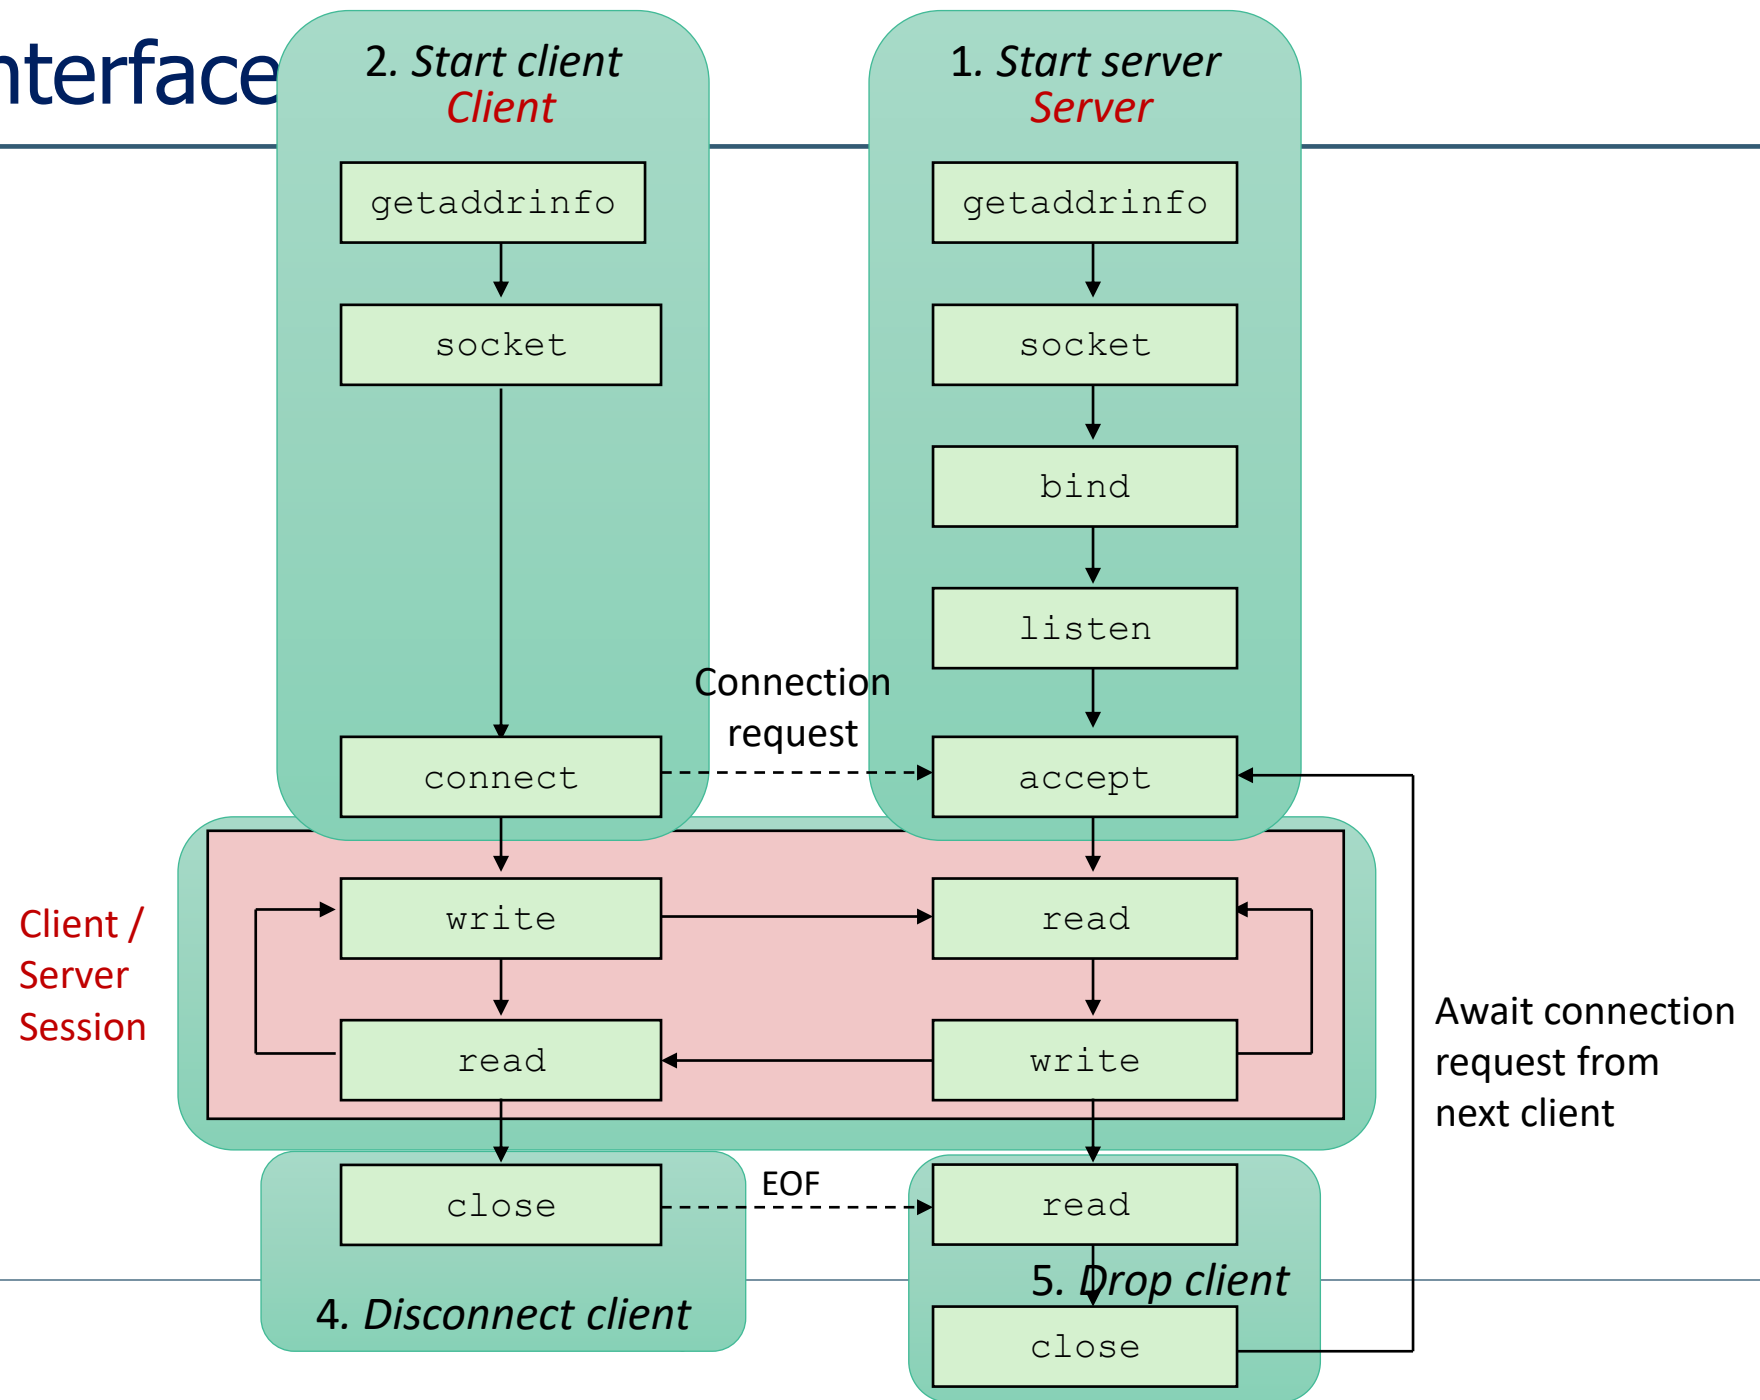
Programmer's View



- Both client and server use socket interface!
- Not only for TCP, but also for UDP

TCP? UDP?

TCP Socket Interface



System Calls for Basic Server

- `getaddrinfo()` & `freeaddrinfo()`
 - Helper function for IP lookup
 - You may use `gethostbyname()` or `gethostbyname_r()`
- `socket()`
 - Makes a socket and allocates system resources for the socket
- `bind()`
 - Binds IP address and port to the socket
- `listen()`
 - Creates a connection queue to allow connections from clients
- `accept()`
 - Retrieves a connection from the connection queue and **returns a socket** for an established connection
- `read()` & `write()`
 - Next slide
- `close()`
 - Close the socket and cleans up resources

For this assignment, use 0.0.0.0
for binding server IP address

Socket descriptor,
Socket structure for metadata,
`sk_buff`, ...

Listen, accept: for TCP
TCP server has 2 sockets

System Calls for Basic Client

- `getaddrinfo()` & `freeaddrinfo()`
 - Helper function for IP lookup
 - You may use `gethostbyname()` or `gethostbyname_r()`
 - `socket()`
 - Makes a socket and allocates system resources for the socket
 - `bind()`
 - Binds IP address and port to the socket
 - `connect()`
 - Creates a connection to the specified address:port (starts TCP handshake)
 - `read()` & `write()`
 - Next slide
 - `close()`
 - Close the socket and cleans up resources
- Socket descriptor,
Socket structure for metadata,
sk_buff, ...
- connect: for TCP
TCP client has 1 socket

read()

- `ssize_t read(int fd, void *buf, size_t nbytes)`
- Returns
 - # of bytes read (> 0)
 - -1 (error, you should check `errno`)
 - No bytes read
 - 0 (closed by the peer)
- `errno == EAGAIN, EWOULDBLOCK`
 - Nothing to read from the TCP socket recv buffer, try later
- `errno == EINTR`
 - Failed due to interrupt, try again right now
- `errno == ECONNRESET`
 - The peer abruptly reset the connection

write()

- `ssize_t write(int fd, void *buf, size_t nbytes)`
- Returns
 - # of bytes wrote (> 0)
 - -1 (error, you should check `errno`)
 - No bytes wrote
- `errno == EAGAIN, EWOULDBLOCK`
 - TCP socket send buffer is full, try later
- `errno == EINTR`
 - Failed due to interrupt, try again right now
- `errno == ECONNRESET`
 - The peer abruptly reset the connection
- `errno == EPIPE`
 - The peer closed the connection

Background 2: Lock

Spin Lock vs. Mutex Lock

- Spin Lock
 - Continuously checks if the lock is available while consuming CPU cycles (busy-waiting)
 - Optimized for short wait times and multi-core environments
- Mutex Lock
 - Puts the waiting thread to sleep, releasing the CPU until the lock becomes available
 - Suitable for longer wait times or resource-intensive applications

Aspect	Spin Lock	Mutex Lock
Waiting	Busy-waiting (CPU is fully utilized)	Sleep-waiting (CPU is released)
Overhead	Low (no context switching)	High (context switching involved)
Use case	Short wait times	Long wait times
Multi-core	Effective in multi-core systems	Works well in both single/multi-core
Complexity	Simple	Relies on OS support

Mutex Lock

- Protects shared data
- Allows only one thread to access the critical section

```
int shared_data = 0;
void *thread_function(void* arg) {
    pthread_mutex_lock(&mutex); // acquire mutex lock

    // critical section: access to the shared data
    shared_data++;
    printf("Thread %ld incremented data: %d\n",
        (long)arg, shared_data);

    pthread_mutex_unlock(&mutex); // release mutex lock
    return NULL;
}
```

- What if most of the accesses are reads?

Many Readers & Few Writers

- Contention necessary?

```
int shared_data = 0;
void *reader_function(void* arg) {
    pthread_mutex_lock(&mutex);

    printf("Reader %ld read data: %d\n",
           (long)arg, shared_data);

    pthread_mutex_unlock(&mutex);
    return NULL;
}
```

```
void *writer_function(void* arg) {
    pthread_mutex_lock(&mutex);

    shared_data++;
    printf("Writer %ld updated data: %d\n",
           (long)arg, shared_data);

    pthread_mutex_unlock(&mutex);
    return NULL;
}
```

- If there is no concurrent writer, multiple threads can concurrently read

Reader-Writer Lock

- A specialized lock that distinguishes b/w read and write access
 1. Allows concurrent multiple readers
 - Blocks writer when any readers exist
 - When there is no more readers, wake up one pending writer (Which one?)
 2. Allows single writer only
 - Blocks other threads when writing
 - When there is no more writers, wake up all pending readers, then wake up one pending writer

Aspect	Mutex Lock	Reader-Writer Lock
Concurrency	Only one thread (reader or writer) at a time	Multiple readers or a single writer
Performance	Suitable for low read-to-write ratio	Optimized for high read-to-write ratio
Complexity	Simpler	More complex
Starvation Risk	No starvation (single queue)	Writer starvation
Use Case	Any critical section	Scenarios with frequent reads and rare writes

- By default, reader-writer lock is usually **reader-priority**

Reader-Priority vs. Writer-Priority (Out of Scope)

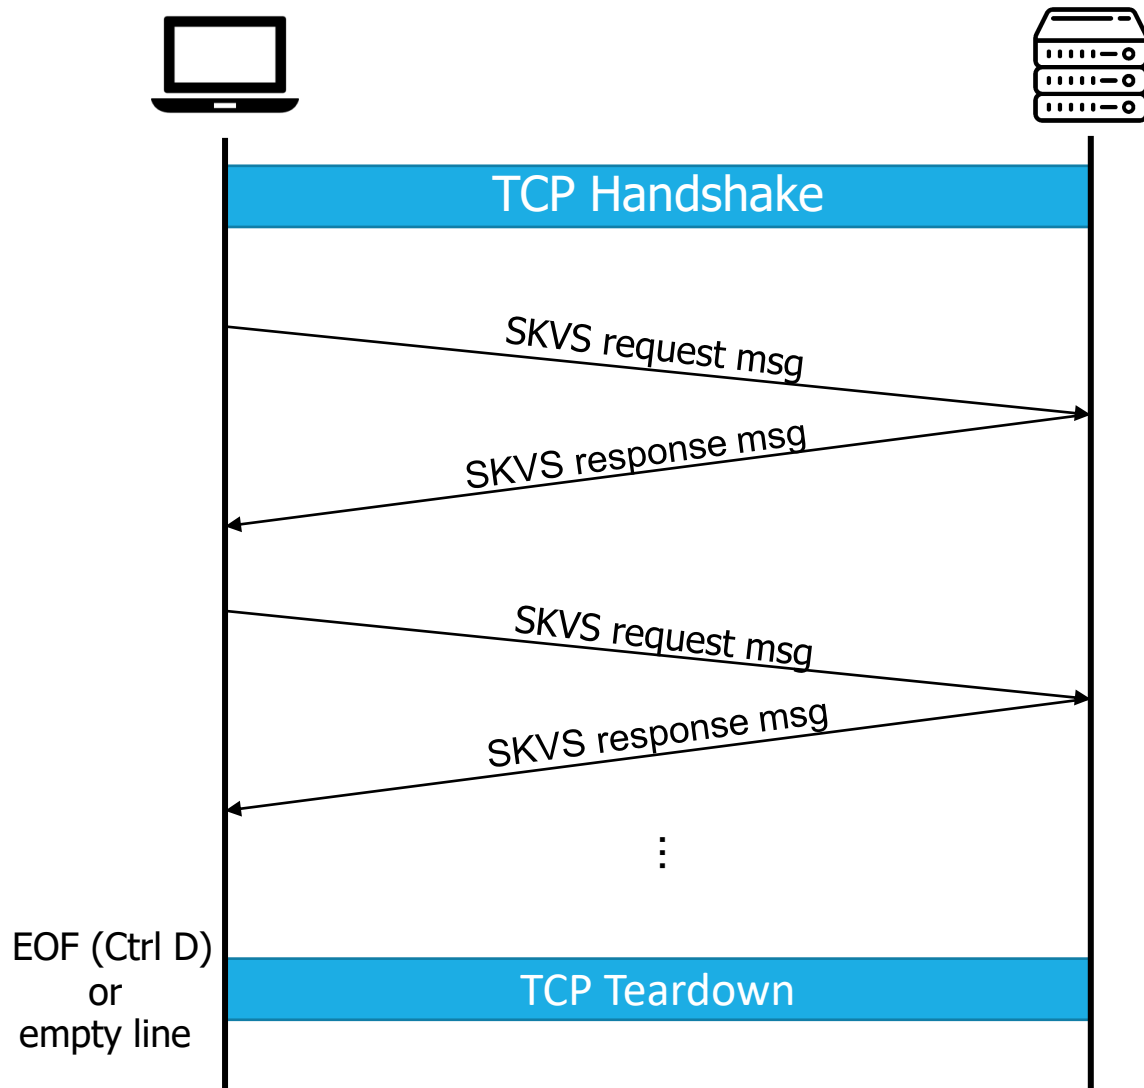
- Reader-Priority RW Lock
 - A read request is granted immediately
 - A write request waits until all ongoing read operations complete
 - If new requests keep arriving, the write request can be indefinitely delayed (writer starvation)
- Writer-Priority RW Lock
 - If a write request is waiting, any new read requests are delayed
 - Once ongoing read operations finish, the write request is executed immediately
 - After the write completes, pending read requests are processed
 - Subsequent write requests continue to take precedence over new reads

Aspect	Reader-Priority	Writer-Priority
Priority	Readers are prioritized	Writers are prioritized
Starvation Risk	Writers may starve	Readers may starve
Performance	Optimized for read-heavy workloads	Balances performance for both reads and writes
Complexity	Simpler	More complex
Use Case	Best for frequent reads	Suitable for fair scheduling b/w reads and writes

Part 1:

Simple Key-Value Store Protocol

SKVS Protocol



- One connection per one client
- **Keep alive** until typing empty line (`\n`) or EOF (Ctrl D)
- **Half-duplex**
 - After sending request, wait for the response
- **Text** based protocol
 - Key and value should be also text
 - Refer to the next slide
- Server should be **stateful**
 - Key-value pairs should be accessible by other clients
- Default service port: 8080

SKVS Protocol (cont.)

- Request Msg Format:

1. [CMD]_s[key]_s[value]_n
 2. [CMD]_s[key]_n
- "_s": a space character
 - "_n": a newline character

- Examples

- "CREATE hello world□_n"
- "READ hello□_n"
- "UPDATE hello snu□_n"
- "DELETE hello□_n"

- CMD: one of CREATE, READ, UPDATE, and DELETE

- case-insensitive (e.g., rEAd, cReAtE are okay)

- Key, value: string without _s nor _n

- No binary, only text, case-sensitive
- len(key) ≤ 32B, len(SKVS msg) ≤ 4096B
 - Including _n

- Response Msg Format:

- CREATE_sOK_n
- [value]_n
- UPDATE_sOK_n
- DELETE_sOK_n
- COLLISION_n
- NOT_sFOUND_n
- INVALID_sCMD_n
- INTERNAL_sERR_n

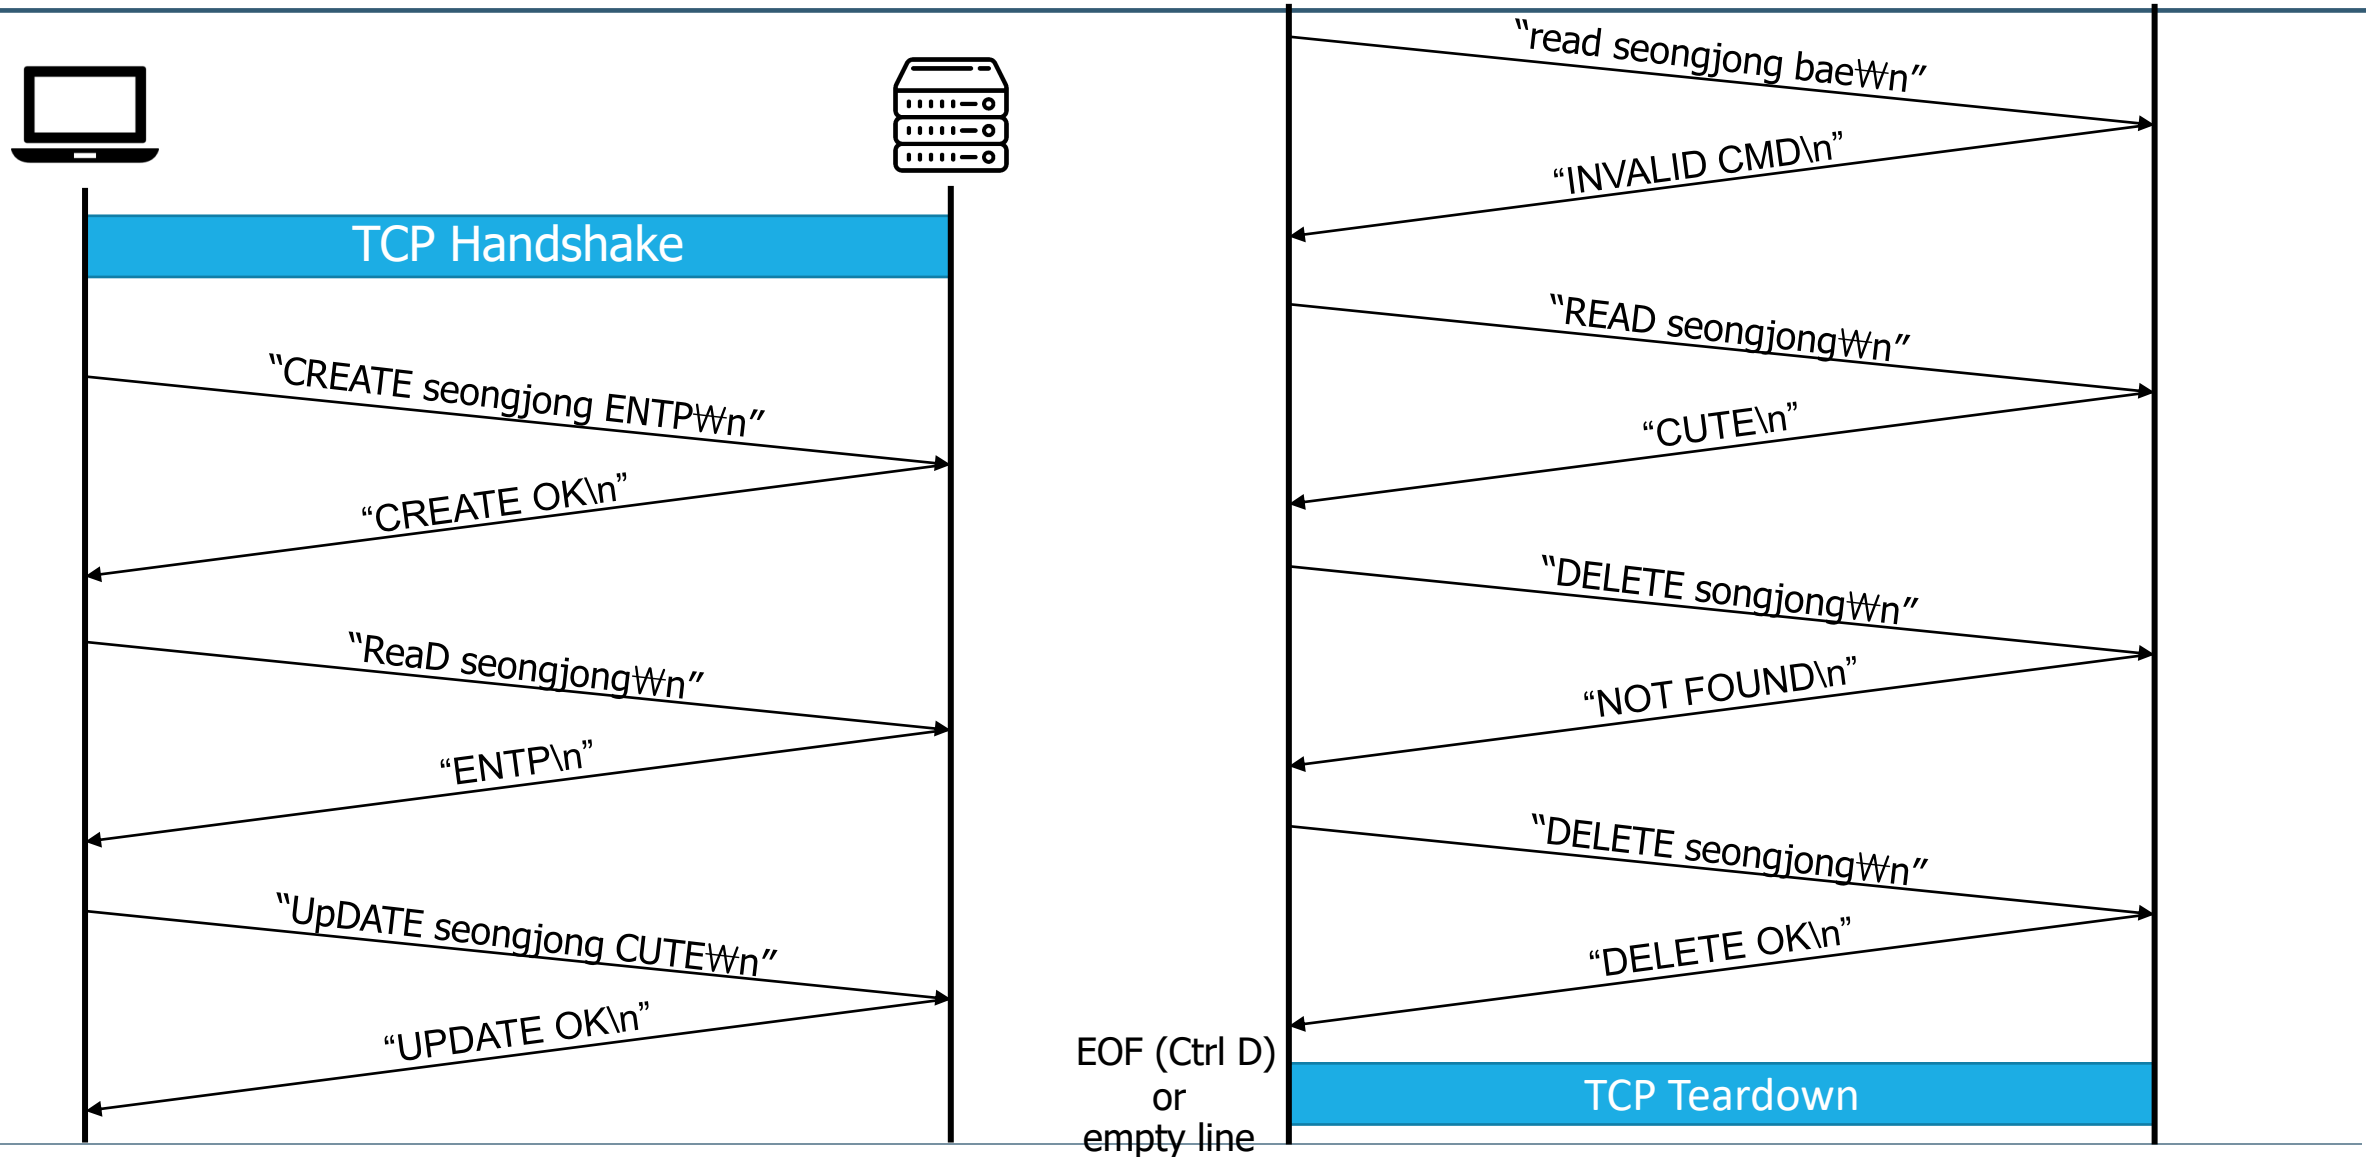
- EOF or empty line on client

- Close the connection
- Exit client program

- Support for file option

- `./client -f input.txt`
- Read the input file as line by line
- Close and exit after receiving last response
- Again, EOF should close and exit

Example

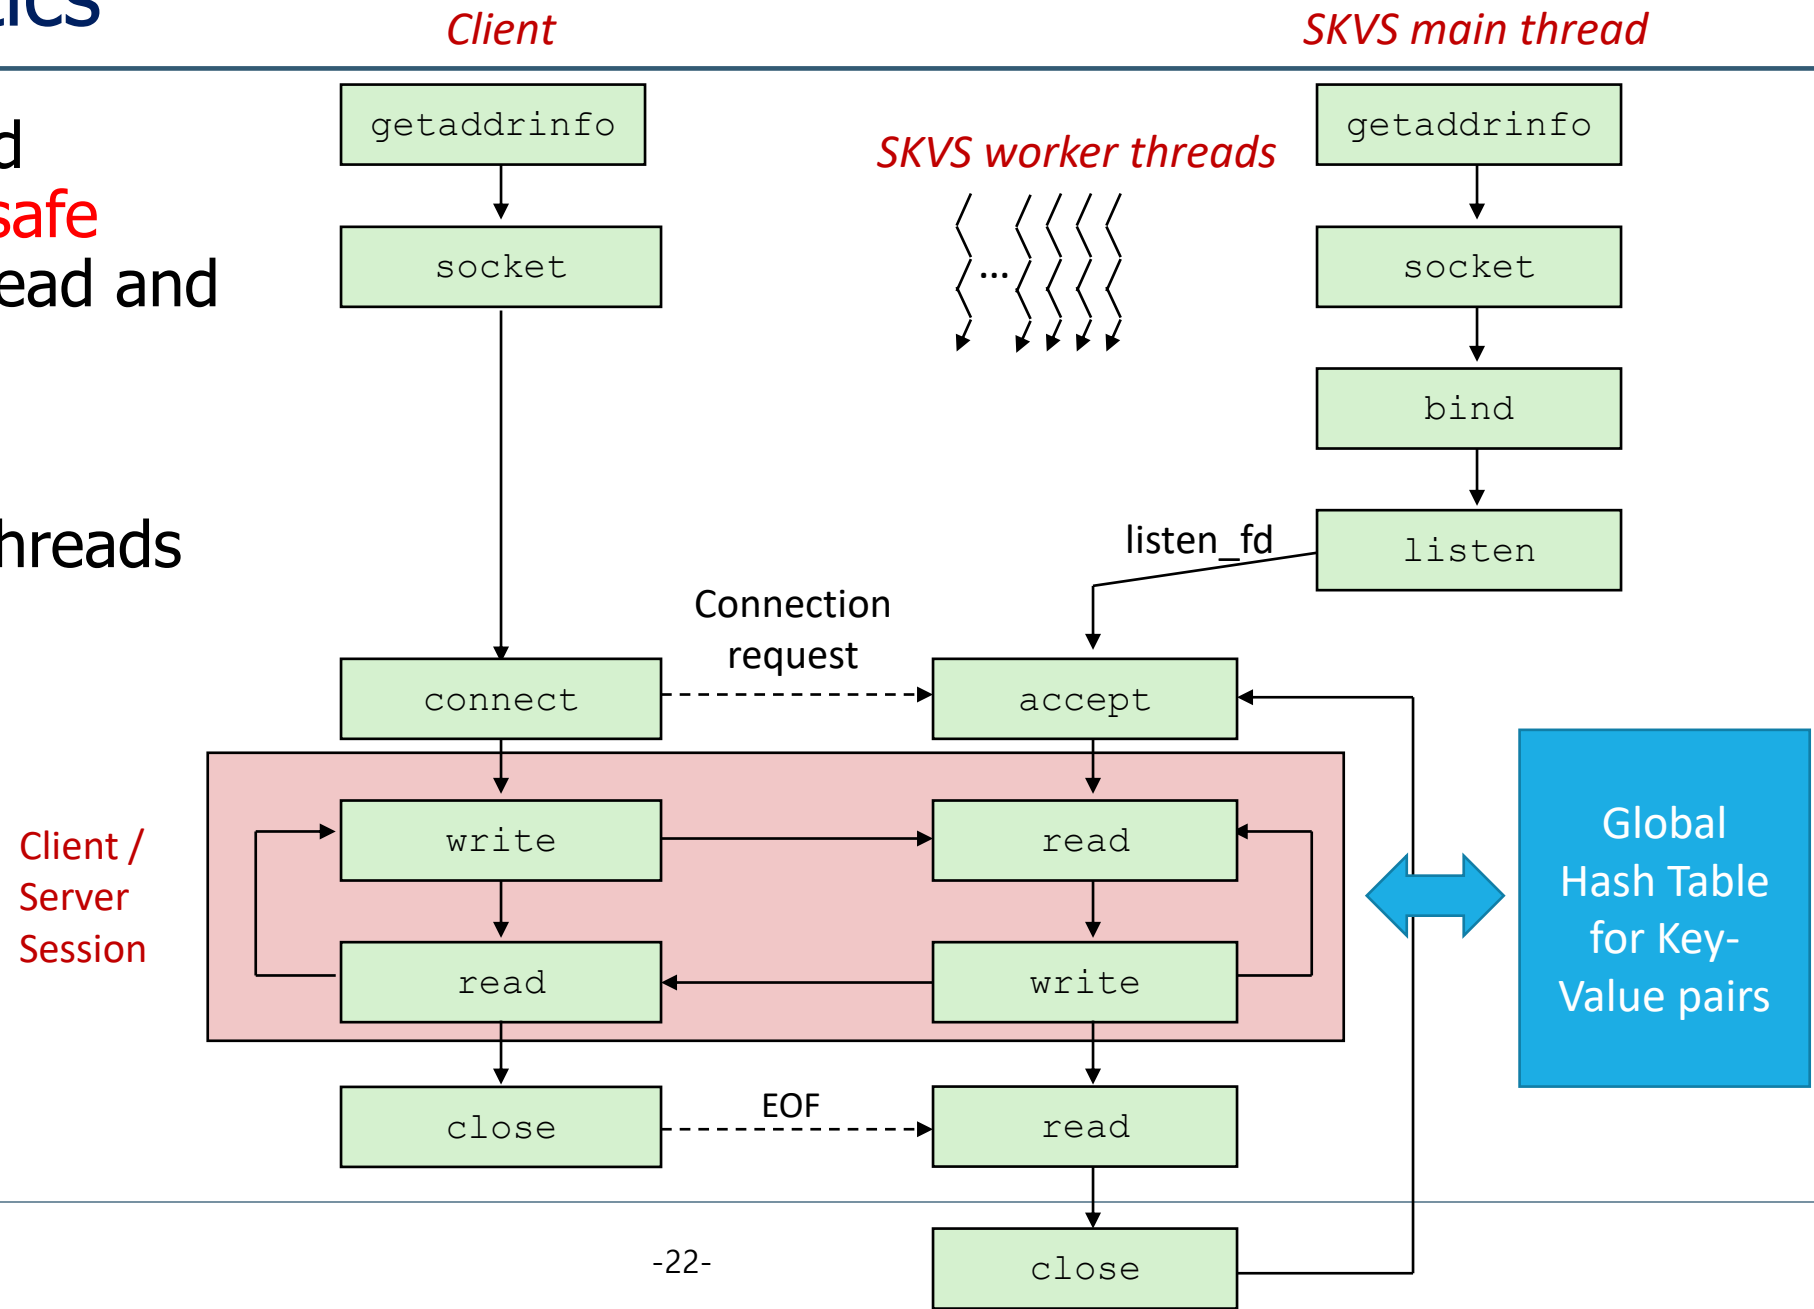


Handling Errors

- If SKVS request is in invalid format?
 - Responds "INVALID CMD□n"
- If the length of SKVS msg > 4096
 - Responds "INVALID CMD□n"
- If the length of [key] > 32
 - Responds "INVALID CMD□n"
- Received CREATE msg, but key-value pair already exists?
 - Responds "COLLISION□n"
- Received READ/UPDATE/DELETE msg, but no such key exists?
 - Responds "NOT FOUND□n"
- If any internal error occurs, so cannot serve the request?
 - Responds "INTERNAL ERR□n"

SKVS Semantics

- Global HT should support **thread-safe access** to both read and write operations
- Use **10** worker threads



Why Static 10 Threads? Any Other Ways for Concurrency?

- What you have learned
 - `accept()`, `read()`, and `write()` are blocking..
 - What if dynamically `fork()` / `pthread_create()` for every clients?
 - (+) Easy to implement
 - (+) Easy to leverage multiple cores
- Reality..
 - (--) Wastes system resources
 - (--) Poor performance due to creating processes or threads
 - (--) Poor performance due to context switching
 - Using `fork()` / `pthread_create()` for each request is inefficient
(`fork()`/`pthread_create()` is heavyweight)

Event-Driven Socket Programming (Out of Scope)

- Set the sockets non-blocking
 - Use event-driven API (e.g., epoll)
 - Static threads to leverage multiple cores
 - Then,
 - (+) Best performance
 - Supports concurrency without context switching
 - Small number of threads
 - (--) Hard to implement using event-driven APIs
- Workaround: Simply use static threads for this assignment! 😊

SKVS API

- Helper API for parsing and processing the SKVS requests
- `struct skvs_ctx *skvs_init(size_t hash_size, int delay)`
 - Initiates SKVS context
- `void skvs_destroy(struct skvs_ctx *ctx)`
 - Destroys SKVS context
- `const char *skvs_serve(struct skvs_ctx *ctx, char *buffer, size_t len)`
 - Serves an appropriate command for the given SKVS request
 - e.g., [value], "INVALID CMD", "CREATE OK", "UPDATE OK", "COLLISION", ...
 - Notice: You should add a line feed at the end

Client & Server Requirement

- Your client should connect **remote server** using domain name with arbitrary port
 - e.g., client on sp03.snucse.org → server on sp04.snucse.org:8080
 - `./server [-p port (8080)] [-t num_threads (10)] [-d rwlock_delay (0)] [-s hash_size (1024)]`
 - `./client [-i server_ip_or_domain (127.0.0.1)] [-p port (8080)] [-t]`
 - Your server should serve **~10** concurrent clients
 - Your server and client should **interoperate with reference client and server, respectively**
- Your server should serve large key-value pair (~32B key, ~4096B SKVS msg)
- Reference server/client binaries will be provided
 - No support for Mac, windows
- For usage, refer to the reference server/client

Usage Example for Interactive Mode (client w/ -t)

- Interactive mode for your better understanding and testing
 - Try it on using our reference client
- We will not use this mode for grading
 - No deduction even if your client does not support -t

```
junghan@sp01:~/lab-5-simple-kvs/assign5/ref$ ./server -p 8000 -t 5 -s 4096
Server listening on 0.0.0.0:8000
0th worker ready
1th worker ready
3th worker ready
2th worker ready
4th worker ready

junghan@sp02:~/lab-5-simple-kvs/assign5/ref$ ./client -i sp01.snucse.org -p 8000 -t
Connected to sp01.snucse.org:8000
Enter command: create hello world
Server reply: CREATE OK
Enter command: read hello
Server reply: world
```

Usage Example for Silent Mode (client w/o -t)

- We will use this mode for grading
 - On the client, SKVS responses should be printed out to stdout
 - No any other messages allowed on the client

```
junghan@sp01:~/lab-5-simple-kvs/assign5/ref$ ./server -p 8000 -t 5 -s 4096
Server listening on 0.0.0.0:8000
1th worker ready
0th worker ready
3th worker ready
2th worker ready
4th worker ready

junghan@sp02:~/lab-5-simple-kvs/assign5/ref$ cat input
create hello world
read hello
junghan@sp02:~/lab-5-simple-kvs/assign5/ref$ ./client -i sp01.snucse.org -p 8000 < input
CREATE OK
world
```

Part 2: Global Hash Table and rwlock

- SKVS library depends on hash table functions
- Hash table functions depend on rwlock functions
- Currently `hashtable.c`, `rwlock.c` are empty, but you should fill them
- `int hash_insert(hashtable_t *table, char *key, char *value)`
 - Needs a **write lock**
 - Duplicates key and value
 - Fills node structure and inserts node to the table
 - Returns 0 if collision, 1 if inserted, -1 if any internal error
- `int hash_search(hashtable_t *table, char *key, char **value)`
 - Needs a **read lock**
 - Returns 0 if not found, 1 (and outputs to `value`) if found, -1 if any internal error

hashtable.c

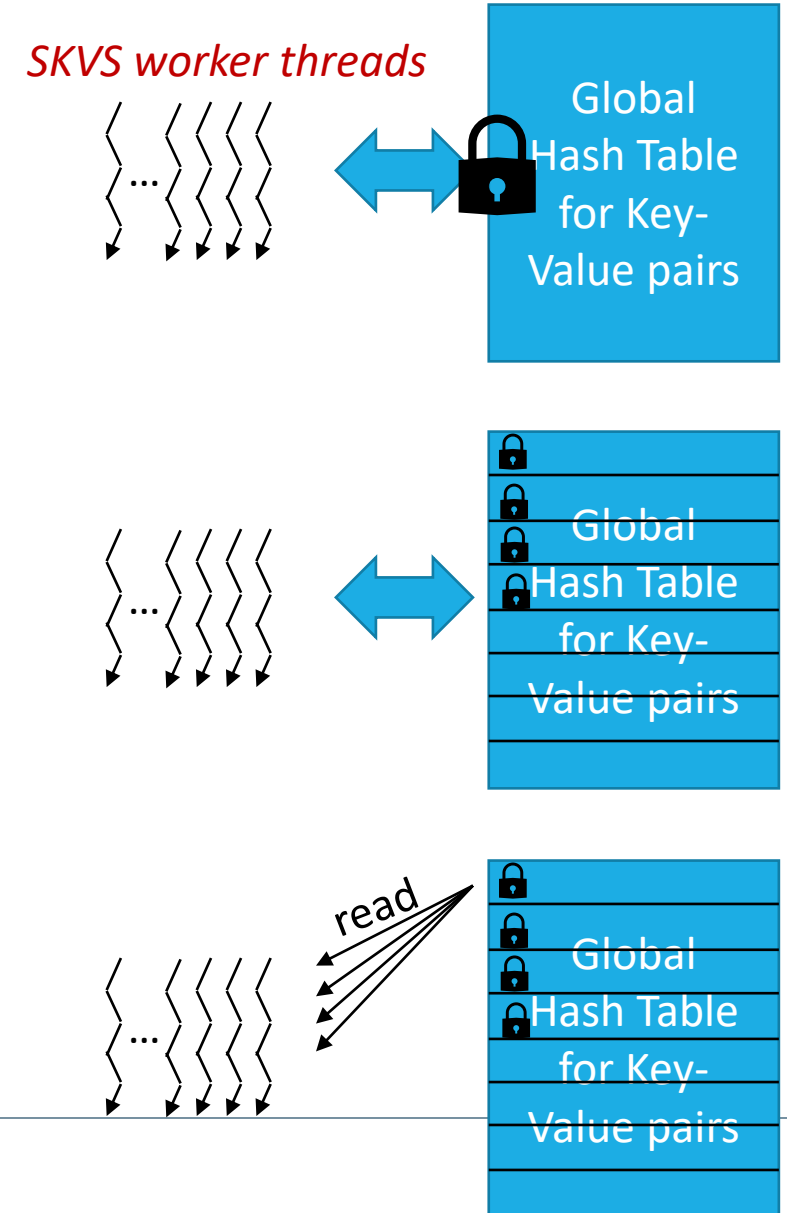
- `int hash_update(hashtable_t *table, char *key, char *value)`
 - Needs a **write lock**
 - Duplicates value
 - Updates node->value to new one
 - Free old value
 - Returns 0 if not found, 1 if updated, -1 if any internal error
- `int hash_delete(hashtable_t *table, char *key)`
 - Needs a **write lock**
 - Frees key and value
 - Evicts from the table
 - Returns 0 if not found, 1 if deleted, -1 if any internal error

Dumping Global Hash Table

- Your SKVS server should dump the hash table using `hash_dump()` before exit on SIGINT
- Register SIGINT handler for dump

Lock?

- SKVS uses global hashtable
 - Any threads consistently access the key-value pair
 - Accesses to the hashtable needs mutex lock
 - But, lock contention leads to poor performance..
1. Fine-grained lock
 - **For each bucket**
 - Less contention compared to course-grained lock
 2. Reader-writer lock (reader-priority)
 - Assumes many readers, few writers
 - **Multiple readers are allowed at the same time**
 - Wakes up readers first, then the **oldest writer**



- Reader-priority reader-writer lock
- **Not** allowed to use `pthread_rwlock` API
- Implement your **own** one in `rwlock.c` using `pthread_mutex` and `pthread_cond` API

- `int rwlock_init(rwlock_t *rw, int delay)`
- `int rwlock_read_lock(rwlock_t *rw)`
- `int rwlock_read_unlock(rwlock_t *rw)`
- `int rwlock_write_lock(rwlock_t *rw)`
- `int rwlock_write_unlock(rwlock_t *rw)`
- `int rwlock_destroy(rwlock_t *rw)`

Wake Up using Condition Variables

```
typedef struct {  
    int read_count;           // number of current readers  
    int write_count;          // number of current writers  
    pthread_mutex_t lock;     // mutex lock for protecting other fields  
    pthread_cond_t readers;   // condvar for threads waiting read  
    pthread_cond_t writers;   // condvar for threads waiting write  
    ...  
} rwlock_t
```

- Example for waking up other threads using condition variables

Thread 1: `pthread_cond_wait(&condvar, &mutex);`

Thread 2: `pthread_cond_signal(&condvar);`

- You should distinguish reader threads and writer threads

rwlock_read_lock

```
int rwlock_read_lock(rwlock_t *rw)
```

1. Acquire mutex lock
2. Wait if any threads are writing
 - By waiting for reader condvar readers
- Wake up!
3. Increment read_count
4. Release mutex lock


```
typedef struct {  
    int read_count;  
    int write_count;  
    pthread_mutex_t lock;  
    pthread_cond_t readers;  
    pthread_cond_t writers;  
    ...  
} rwlock_t
```

rwlock_write_lock

```
int rwlock_write_lock(rwlock_t *rw)
```

1. Acquire mutex lock
2. Insert this thread to `writer_ring`
3. Wait if any threads are writing
 - By waiting for writer condvar `writers`
- Wake up!
4. Check if this thread is the oldest
 - If not, wait again
5. Increment `write_count`
6. Evict this thread from `writer_ring`
7. Release mutex lock

```
typedef struct {  
    int read_count;  
    int write_count;  
    pthread_mutex_t lock;  
    pthread_cond_t readers;  
    pthread_cond_t writers;  
  
    // writer ring  
    pthread_t *writer_ring;  
    int writer_ring_head;  
    int writer_ring_tail;  
} rwlock_t
```



How to Test Write Lock?

- Before testing, check below
 1. Implement server/client first with **robust read/write**
 2. Your server **must** be ready for concurrency first
 3. `./server -d 5`
 - Add `sleep(5)` in `rwlock_write_unlock()`
 - Add `sleep(5)` in `rwlock_read_unlock()`
- Example test scenario
 1. Send "CREATE hello 1□n" on client 1
 - It will hold a write lock for 10 seconds
 2. Send "CREATE hello 2□n" on client 2
 - Check whether client 2 gets stuck
 3. Send "CREATE hello 3□n" on client 3
 - Check whether client 3 gets stuck
 4. Send "CREATE hello 4□n" on client 4
 - Check whether client 4 gets stuck
 7. After 5 seconds...
 - Check whether client 1 receives "CREATE OK□n"
 8. After 5 seconds...
 - Check whether client 2 receives "COLLISION□n"
 9. After 5 seconds...
 - Check whether client 3 receives "COLLISION□n"
 10. After 5 seconds...
 - Check whether client 4 receives "COLLISION□n"

How to Test Read Lock?

- Before testing, check below
 1. Implement server/client first with **robust read/write**
 2. Your server **must** be ready for concurrency first
 3. `./server -d 5`
 - Add `sleep(5)` in `rwlock_write_unlock()`
 - Add `sleep(5)` in `rwlock_read_unlock()`
- Example test scenario
 1. Send "CREATE hello world□n" on client 1
 - It will hold a write lock for 10 seconds
 2. Send "READ hello□n" on client 2
 3. Send "READ hello□n" on client 3
 4. Send "READ hello□n" on client 4

simultaneously

 7. After 5 seconds...
 - Check whether client 1 receives "CREATE OK□n"
 8. After 5 seconds...
 - Check whether client 2 receives "world□n"
 - Check whether client 3 receives "world□n"
 - Check whether client 4 receives "world□n"

Complex Scenario

- Before testing, check below

1. Implement server/client first with **robust read/write**
2. Your server **must** be ready for concurrency first
3. `./server -d 5`
 - Add `sleep(5)` in `rwlock_write_unlock()`
 - Add `sleep(5)` in `rwlock_read_unlock()`

- Example test scenario

1. Send "CREATE hello world\n" on client 1
 - It will hold a write lock for 10 seconds
2. Send "DELETE bye\n" on client 2
 - Check whether client 2 receives "NOT FOUND\n" immediately (why?)
3. Send "READ hello\n" on client 2
 - Check whether client 2 gets stuck
4. Send "UPDATE hello snu\n" on client 3
 - Check whether client 3 gets stuck
5. Send "DELETE hello\n" on client 4
 - Check whether client 4 gets stuck
6. Send "READ hello\n" on client 5
 - Check whether client 5 gets stuck

7. After 5 seconds...

- Check whether client 1 receives "CREATE OK\n"
- Check whether client 2 receives "world\n"
- Check whether client 5 receives "world\n"

8. After 5 seconds...

- Check whether client 3 receives "UPDATE OK\n"

9. After 5 seconds...

- Check whether client 4 receives "DELETE OK\n"

Requirements Summary

- Your client should connect remote server using domain name with arbitrary port
 - e.g., client on sp03.snucse.org → server on sp04.snucse.org:8080
 - ```
./server [-p port (8080)] [-t num_threads (10)] [-d rwlock_delay (0)]
[-s hash_size (1024)]
```
  - ```
./client [-i server_ip_or_domain (127.0.0.1)] [-p port (8080)] [-t]
```
 - Your server should serve ~10 concurrent clients
 - Your server and client should interoperate with reference client and server, respectively
- Your server should serve large key-value pair (~32B key, ~4096B SKVS msg)
- Print the entries in the global hash table using `hash_dump()`
 - Do not modify this function
- Concurrent multiple readers' access to the global hash table
- Correct reader-priority RW lock semantic
- **Do not modify other files than** `client.c`, `server.c`, `hashtable.c`, and `rwlock.c`

Guidelines

Notice

- Do not change the name and the prototype of the skeleton code
- What to submit
 - A tarball named `202412345_assign5` including `server.c` `client.c` `hashtable.c` `rwlock.c`.

```
cd assign5/src  
make submit ID=202412345
```

- Replace `202412345` to your student ID without dash

Deadline

- **Deadline: ~ 2024. 12. 20 21:00**
 - 0 Points if deadline is missed
 - 0 Points for copying
- **Contact**
 - Lab 5 TA e-mail: cerotyki@snu.ac.kr
 - TA mailing list: snu-sysp@googlegroups.com

Reference Binaries for self-checking

- Reference server/client binaries will be provided
 - No support for Mac, windows
- For any ambiguities, refer to the reference server/client
- You may use `TRACE_PRINT()` and `DEBUG_PRINT()` for debugging

```
CFLAGS += -DTRACE
CFLAGS += -DDEBUG
```

Q&A